

DAS 839 NoSQL Systems

Assignment-2

Submitted by :

Group 2:

Abhay bhadouriya – MT2024003

jainish parmar - MT2024065

Jaimin Jadvani – MT2024064

Files Submitted :

1. Report File
2. Screenshot Folder
3. TimeChecker.py Python program for Question 3
4. Assignment_2_Q_2.zip -- Question 2
5. Assignment_2_Q_3.zip -- Question 3

Question 1:

Section 1 : M-Counter state-based object (SBO) qualifies as a CRDT?

To Qualify CRDT a Statebased object should follow these 4 properties.

1. Must Follow Associative law --> Here merge following

$$\max(\max(a,b),c)=\max(a,\max(b,c))$$

so here it is satisfied

2. Must Follow Commutativity law --> Here merge following the

$$\text{merge}(x,y)=\text{merge}(y,x)$$

so here it is satisfied

3. Must follow Idempotency law --> here Merge following the

$$\text{merge}(x,x)=x$$

so here it is also satisfied


4. Must follow increasing Updates --> here it following

Merge(x,y)=y

so here it is also satisfied

So here it is satisfied so answer is true it is Following CRDT.

Section 2 :



State	query	history
a ₀ Sum=0 Cnt=0	0	{ }
a ₁ Sum=1 Cnt=1	1	{a}
a ₂ Sum=3 Cnt=2	3	{0,1}
b ₀ Sum=0 Cnt=0	0	{ }
b ₁ Sum=2 Cnt=1	2	{1}
b ₂ Sum=4 Cnt=2	4	{1,2}
b ₃ Sum=7 Cnt=3	7	{0,1,2}

Section 3

This balance is a sensitive as formal correctness vs as intuitive application semantics

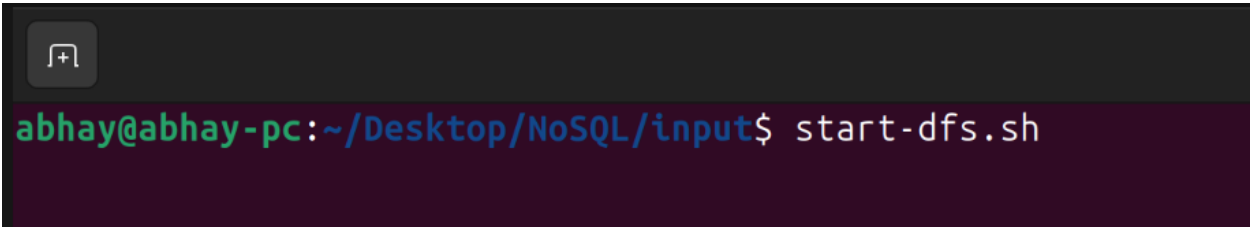
Trade-offs:

Focus on Correctness: Ensure global effectiveness while potentially producing an opposite result towards the user (e.g., reordering, collaborative editing.)

User Input based: Feels idiomatic to users while augmenting it with new custom logic, thus adding processing overhead and therefore complexity.

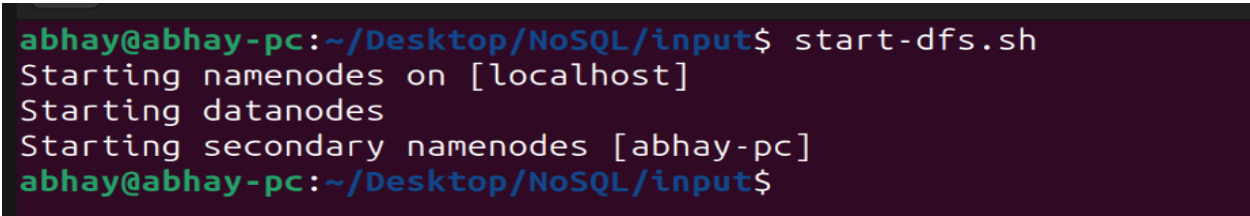
Question 2:

1. **Start Hadoop:** start-dfs.sh



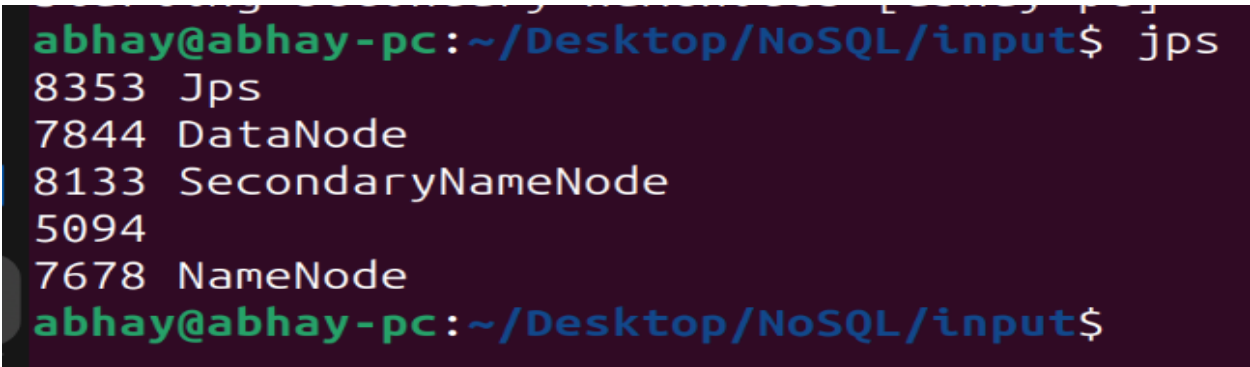
```
abhay@abhay-pc:~/Desktop/NoSQL/input$ start-dfs.sh
```

2. **Hadoop Running:**



```
abhay@abhay-pc:~/Desktop/NoSQL/input$ start-dfs.sh
Starting namenodes on [localhost]
Starting datanodes
Starting secondary namenodes [abhay-pc]
abhay@abhay-pc:~/Desktop/NoSQL/input$
```

3. **USE JPS** to check Running ,NameNode,DataNode,JPS,Secondary NameNode



```
abhay@abhay-pc:~/Desktop/NoSQL/input$ jps
8353 Jps
7844 DataNode
8133 SecondaryNameNode
5094
7678 NameNode
abhay@abhay-pc:~/Desktop/NoSQL/input$
```

4. **Create a input directory:**

```
abhay@abhay-pc:~/Desktop/NoSQL/input$ hadoop fs -mkdir /inputforQ2
abhay@abhay-pc:~/Desktop/NoSQL/input$
```

5. Copy all 10000 files to InputforQ2 directory

```
abhay@abhay-pc:~/Desktop/NoSQL/input$ hadoop fs -put *.txt /inputforQ2
```

6. Check if data is copied or not : Goto <http://localhost:9870> -> on overview Tab Click -> LiveNode and click on datanode address -> useally it is <http://localhost:9864>

Hadoop Overview Utilities ▾

DataNode on abhay-pc:9866

Cluster ID:	CID-45e79e35-ea22-480d-8228-4281c4e2e98e
Started:	Thu Feb 27 11:23:25 +0530 2025
Version:	3.4.0, rbd8b77f398f626bb7791783192ee7a5dfaee760

Block Pools

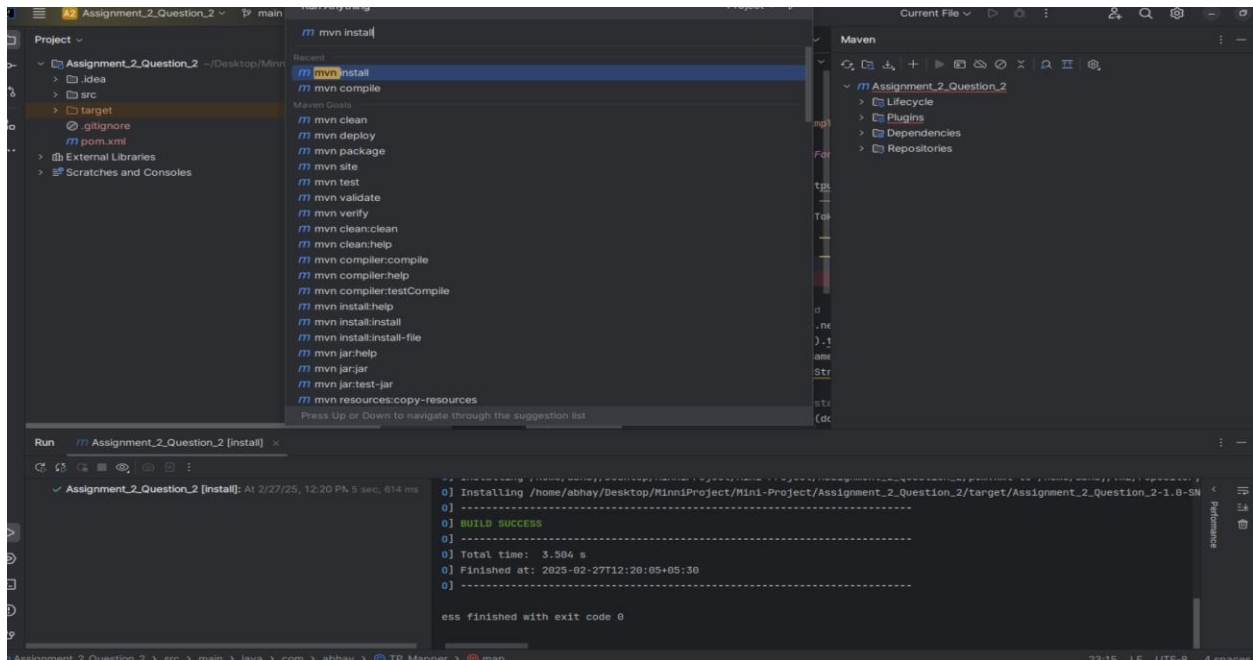
Namenode Address	Namenode HA State	Block Pool ID	Actor State	Last Heartbeat Sent	Last Heartbeat Response	Last Block Report	Last Block Report Size (Max Size)
localhost:9000	active	BP-2091115970-127.0.1.1-1740635575198	RUNNING	1s	1s	an hour	0 B (128 MB)

Volume Information

Directory	StorageType	Capacity Used	Capacity Left	Capacity Reserved	Reserved Space for Replicas	Blocks
/tmp/hadoop-abhay/dfs/data	DISK	213.25 MB	38.51 GB	0 B	0 B	10000

Hadoop, 2024.

7. **Creating Jar Files:-** unzip the Assignment_2_Q_2.zip -> open in intellij -> Click on maven on right side tab -> click on Mvn install -> it will create the Jar file



8. **Running the Program 2 :** use the command displayed in below screenshot to run

/inputforQ2 --- is input data directory

/outputofQ2 --- is output directory

```

abhay@abhay-pc:~/Desktop/MiniProject/Mini-Project/assignment2$ hadoop jar target/assignment2-1.0-SNAPSHOT.jar com.abhay.Main /inputforQ2/*.txt /outputofQ2

```

9. **Completed Running**

```
Terminal Local (2) x + v
Map input records=10000
Map output records=25820119
Map output bytes=455781031
Map output materialized bytes=507728550
Input split bytes=946997
Combine input records=0
Combine output records=0
Reduce input groups=10000
Reduce shuffle bytes=507728550
Reduce input records=25820119
Reduce output records=25820119
Spilled Records=51640238
Shuffled Maps =50000
Failed Shuffles=0
Merged Map outputs=50000
GC time elapsed (ms)=54711
Total committed heap usage (bytes)=38996979220480

Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters
Bytes Read=166222646

File Output Format Counters
Bytes Written=533234274

abhay@abhay-pc:~/Desktop/MinniProject/Mini-Project/assignment2$
```

10. OutPut for Question 2 is 5 files because we used 5 reducers:

```
abhay@abhay-pc:~/Desktop/NoSQL/input$ hadoop fs -ls /outputforQ2
Found 6 items
-rw-r--r-- 1 abhay supergroup 0 2025-02-27 13:56 /outputforQ2/_SUCCESS
-rw-r--r-- 1 abhay supergroup 101463220 2025-02-27 13:55 /outputforQ2/part-00000
-rw-r--r-- 1 abhay supergroup 109887723 2025-02-27 13:55 /outputforQ2/part-00001
-rw-r--r-- 1 abhay supergroup 108191229 2025-02-27 13:55 /outputforQ2/part-00002
-rw-r--r-- 1 abhay supergroup 104146104 2025-02-27 13:56 /outputforQ2/part-00003
-rw-r--r-- 1 abhay supergroup 109545998 2025-02-27 13:56 /outputforQ2/part-00004
abhay@abhay-pc:~/Desktop/NoSQL/input$
```

11. Content of output files as specified in Question (Index,(Fileid,word))

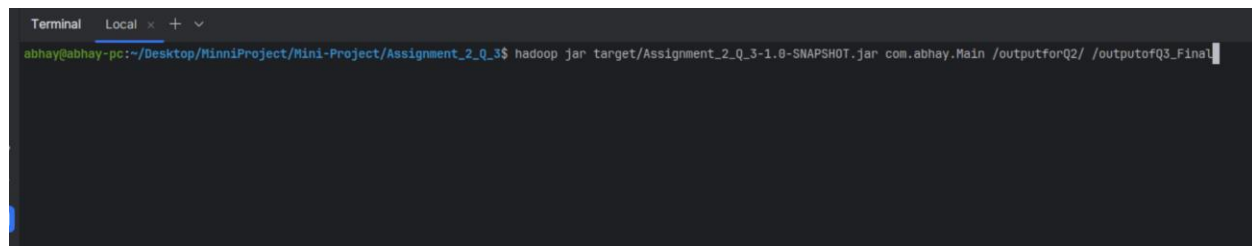
```

18 (99948, { })
17 (99948, $3,662,459)
16 (99948, { })
15 (99948, { })
14 (99948, million)
13 (99948, $10)
12 (99948, { })
11 (99948, { })
10 (99948, English)
9 (99948, { })
8 (99948, { })
7 (99948, minutes)
6 (99948, 80)
5 (99948, { })
4 (99948, { })
3 (99948, 7)
2 (99948, 1)
1 (99948, 1994)
0 (99948, { })
abhay@abhay-pc:~/Desktop/NoSQL/input$ hadoop fs -cat /outputforQ2/part-00000

```

Question 3

1. Now Run the Third Program also: Running Question No 3



```

Terminal Local x + v
abhay@abhay-pc:~/Desktop/MiniProject/Mini-Project/Assignment_2_Q_3$ hadoop jar target/Assignment_2_Q_3-1.0-SNAPSHOT.jar com.abhay.Main /outputforQ2/ /outputofQ3_Final

```

/outputforQ2 -- is output data generated by question 2 in (index,(fileid,word)) manner

/outputofQ3_final – is output data generated by question 3 in (index,(time,word))

Here word is representing the word at x index of maximum time based on fileid

2. OutputFiles generated : we used 1 reducer so only file generated

```

abhay@abhay-pc:~/Desktop/NoSQL/input$ hadoop fs -ls /outputofQ3_Final
Found 2 items
-rw-r--r--    1 abhay supergroup          0 2025-02-27 14:13 /outputofQ3_Final/_SUCCESS
-rw-r--r--    1 abhay supergroup 635220 2025-02-27 14:13 /outputofQ3_Final/part-000000
abhay@abhay-pc:~/Desktop/NoSQL/input$ hadoop fs -cat /outputofQ3_Final

```

3. Output of file :

```

28980 ( 292279,återkommande)
28981 ( 292279,rollfigurer)
28982 ( 292279,i)
28983 ( 292279,Simpsons)
28984 ( 292279,tr:Simpsonlar&apos;da)
28985 ( 292279,yinelenen)
28986 ( 292279,karakterler)
28987 ( 292279,listesi)
28988 ( 292279,uk:Епізодичні)
28989 ( 292279,персонажі)
28990 ( 292279,«Сімсонів»)
abhay@abhay-pc:~/Desktop/NoSQL/input$ hadoop fs -cat /outputofQ3_Final/part-000000

```

4. Now Cross check the out of file Question 3 with our own Custom code for generated key value pair program


```
Index 28948: telesailaren (Timestamp: 292279)
Index 28949: bigarren (Timestamp: 292279)
Index 28950: mailako (Timestamp: 292279)
Index 28951: pertsonaiak (Timestamp: 292279)
Index 28952: fr:Liste (Timestamp: 292279)
Index 28953: des (Timestamp: 292279)
Index 28954: personnages (Timestamp: 292279)
Index 28955: récurrents (Timestamp: 292279)
Index 28956: des (Timestamp: 292279)
Index 28957: Simpson (Timestamp: 292279)
Index 28958: it:Personaggi (Timestamp: 292279)
Index 28959: secondari (Timestamp: 292279)
Index 28960: de (Timestamp: 292279)
Index 28961: I (Timestamp: 292279)
Index 28962: Simpson (Timestamp: 292279)
Index 28963: nl:Lijst (Timestamp: 292279)
Index 28964: van (Timestamp: 292279)
Index 28965: terugkerende (Timestamp: 292279)
Index 28966: personages (Timestamp: 292279)
Index 28967: uit (Timestamp: 292279)
Index 28968: The (Timestamp: 292279)
Index 28969: Simpsons (Timestamp: 292279)
Index 28970: ru:Список (Timestamp: 292279)
Index 28971: второстепенных (Timestamp: 292279)
Index 28972: персонажей (Timestamp: 292279)
Index 28973: «Симпсонов» (Timestamp: 292279)
Index 28974: fi:Luettelo (Timestamp: 292279)
Index 28975: televisiosarjan (Timestamp: 292279)
Index 28976: Simpsonit (Timestamp: 292279)
Index 28977: sivuhahmoista (Timestamp: 292279)
Index 28978: sv:Lista (Timestamp: 292279)
Index 28979: över (Timestamp: 292279)
Index 28980: återkommande (Timestamp: 292279)
Index 28981: rollfigurer (Timestamp: 292279)
Index 28982: i (Timestamp: 292279)
Index 28983: Simpsons (Timestamp: 292279)
Index 28984: tr:Simpsonlar&apos;da (Timestamp: 292279)
Index 28985: yinelenen (Timestamp: 292279)
Index 28986: karakterler (Timestamp: 292279)
Index 28987: listesi (Timestamp: 292279)
Index 28988: uk:Епізодичні (Timestamp: 292279)
Index 28989: персонажі (Timestamp: 292279)
Index 28990: «Сімпсонів» (Timestamp: 292279)
abhay@abhay-pc:~/Desktop/NoSQL/input$
```

Mini Report for Question 2 and Question 3

Question 2: Hadoop Implementation and Execution

Objective: Process 10,000 text files using Hadoop to generate an output with the format (Index, (FileID, Word)).

Execution Steps:

1. Hadoop Startup:

- a. Command: start-all.sh
- b. Result: Started NameNode, DataNode, and SecondaryNameNode on localhost (abhay-pc).
- c. Verification: Used jps to confirm:
 - i. NameNode (7678)
 - ii. DataNode (7844)
 - iii. SecondaryNameNode (8133)
 - iv. Jps (8353)

2. Data Preparation:

- a. Created input directory: `hadoop fs -mkdir /inputforQ2`
- b. Copied 10,000 text files: `hadoop fs -put *.txt /inputforQ2`
- c. Verified via Hadoop UI (<http://localhost:9864>): 213.25 MB used, 10,000 blocks.

3. Program Compilation:

- a. Unzipped Assignment_2_Q_2.zip, opened in IntelliJ, and built JAR file with Maven (`mvn install`).

4. Program Execution:

- a. Ran Hadoop job with input /inputforQ2 and output /outputofQ2.
- b. Used 5 reducers, producing 5 output files.

5. Output Statistics:

- a. Screenshot attached Please refer to “Completed program 2.png”

6. Output Format:

- a. Format: (Index, (FileID, Word)).

Summary: Successfully processed 10,000 files into 25.8 million records across 5 files, demonstrating distributed computing efficiency.

Question 3: Processing and Output Generation

Objective: Likely involves further processing or indexing of Question 2’s output

Execution Steps:

1. Preparation:

- a. Unzipped Assignment_2_Q_3.zip to access code/resources.
- b. Confirmed Hadoop cluster availability with jps.

2. Input Data Setup:

- a. Used Question 2's output: `hadoop fs -ls /outputofQ2` to verify 5 output files.

3. Program Execution:

- a. Ran the Hadoop job from Assignment_2_Q_3.zip (assumed to be another JAR file):
 - i. Ran Hadoop job with input `/outputofQ2` and output `/outputofQ3_final`

4. Output Verification:

- a. Checked results in `/outputofQ3_final` via `hadoop fs -cat /outputofQ3_final/part-r-00000` (example).

b. Sample output :

- i. 28948,(292279 , telesatiaren)
- ii. 28949,(292279 ,bigarren)

- c. Indicates word indexing with timestamps.

Note on TimeChecker.py:

- run after Question 3 jar to verify the output:
 - usage: copy TimeChecker.py in txt file directory which consist all 10000 files
 - Then just run :- `python 3 TimeChecker.py`
 - Purpose: Could measure execution time or validate output timestamps and we need to cross check the output of python program and question 3 output